

Perfect secrecy

perfectly secret

- Perfectly secret :Encryption schemes that are provably secure even against an adversary with unbounded computational power.

Probability review

- *Random variable (r.v.):* variable that takes on (discrete) values with certain probabilities
- Probability distribution for a r.v. specifies the probabilities with which the variable takes on each possible value
 - Each probability must be between 0 and 1
 - The probabilities must sum to 1

Probability review

- *Event*: a particular occurrence in some experiment
 - $\Pr[E]$: probability of event E
- Conditional probability: probability that one event occurs, *given that* some other event occurred
 - $\Pr[A \mid B] = \Pr[A \text{ and } B] / \Pr[B]$
- Two random variables X, Y are *independent* if for all x, y : $\Pr[X=x \mid Y=y] = \Pr[X=x]$

Probability review

- Law of total probability: say E_1, \dots, E_n are a *partition* of all possibilities. Then for any A :

$$\Pr[A] = \sum_i \Pr[A \text{ and } E_i] = \sum_i \Pr[A \mid E_i] \cdot \Pr[E_i]$$

Notation

- \mathcal{K} (key space) – set of all possible keys
- \mathcal{C} (ciphertext space) – set of all possible ciphertexts

Probability distributions

- Let M be the random variable denoting the value of the message
 - M ranges over \mathcal{M}
 - Context dependent!
 - Reflects the likelihood of different messages being sent, given the attacker's prior knowledge
 - E.g.,
$$\Pr[M = \text{"attack today"}] = 0.7$$
$$\Pr[M = \text{"don't attack"}] = 0.3$$

Probability distributions

- Let K be a random variable denoting the key
 - K ranges over \mathcal{K}
- Fix some encryption scheme $(\text{Gen}, \text{Enc}, \text{Dec})$
 - Gen defines a probability distribution for K :
$$\Pr[K = k] = \Pr[\text{Gen outputs key } k]$$

Probability distributions

- Random variables M and K are *independent*
 - Require that parties don't pick the key based on the message, or the message based on the key

Probability distributions

- Fix some encryption scheme (Gen, Enc, Dec), and some distribution for M
- Consider the following (randomized) experiment:
 1. Generate a key k using Gen
 2. Choose a message m , according to the given distribution
 3. Compute $c \leftarrow \text{Enc}_k(m)$
- This defines a distribution on the ciphertext!
- Let C be a random variable denoting the value of the ciphertext in this experiment

Example 1

- Consider the shift cipher
 - So for all $k \in \{0, \dots, 25\}$, $\Pr[K = k] = 1/26$
- Say $\Pr[M = 'a'] = 0.7$, $\Pr[M = 'z'] = 0.3$
- What is $\Pr[C = 'b']$?
 - Either $M = 'a'$ and $K = 1$, or $M = 'z'$ and $K = 2$
 - $\Pr[C='b'] = \Pr[M='a'] \cdot \Pr[K=1] + \Pr[M='z'] \cdot \Pr[K=2]$
 $= 0.7 \cdot (1/26) + 0.3 \cdot (1/26)$
 $= 1/26$

Example 1

We can calculate conditional probabilities as well. For example, what is the probability that the message **a** was encrypted, given that we observe ciphertext **B**? Using Bayes' Theorem (Theorem A.8) we have

$$\begin{aligned}\Pr[M = \mathbf{a} \mid C = \mathbf{B}] &= \frac{\Pr[C = \mathbf{B} \mid M = \mathbf{a}] \cdot \Pr[M = \mathbf{a}]}{\Pr[C = \mathbf{B}]} \\ &= \frac{\Pr[C = \mathbf{B} \mid M = \mathbf{a}] \cdot 0.7}{1/26}.\end{aligned}$$

Note that $\Pr[C = \mathbf{B} \mid M = \mathbf{a}] = 1/26$, since if $M = \mathbf{a}$ then the only way $C = \mathbf{B}$ can occur is if $K = 1$ (which occurs with probability $1/26$). We conclude that $\Pr[M = \mathbf{a} \mid C = \mathbf{B}] = 0.7$. \diamond

Example 2

- Consider the shift cipher, and the distribution on M given by

Consider the shift cipher again, but with the following distribution over \mathcal{M} :

$$\Pr[M = \text{kim}] = 0.5, \Pr[M = \text{ann}] = 0.2, \Pr[M = \text{boo}] = 0.3.$$

What is the probability that $C = \text{DQQ}$? The only way this ciphertext can occur is if $M = \text{ann}$ and $K = 3$, or $M = \text{boo}$ and $K = 2$, which happens with probability $0.2 \cdot 1/26 + 0.3 \cdot 1/26 = 1/52$.

We can also compute the probability that `ann` was encrypted, conditioned on observing the ciphertext `DQQ`? A calculation as above using Bayes' Theorem gives $\Pr[M = \text{ann} \mid C = \text{DQQ}] = 0.4$. \diamond

Perfect secrecy (informal)

- “Regardless of any *prior* information the attacker has about the plaintext, the ciphertext should leak no *additional* information about the plaintext”

a scheme to be perfectly secret, observing this ciphertext should have *no effect* on the adversary’s knowledge regarding the actual message that was sent; in other words, the *a posteriori* probability that some message $m \in \mathcal{M}$ was sent, conditioned on the ciphertext that was observed, should be no different from the *a priori* probability that m would be sent. This means that the ciphertext reveals nothing about the underlying plaintext, and the adversary learns absolutely nothing about the plaintext that was encrypted. Formally:

Perfect secrecy (informal)

- Attacker's information about the plaintext = attacker-known *distribution* of M
- Perfect secrecy means that observing the ciphertext should not change the attacker's knowledge about the distribution of M

Perfect secrecy (formal)

- Encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secret* if for every distribution over \mathcal{M} , every $m \in \mathcal{M}$, and every $c \in \mathcal{C}$ with $\Pr[C=c] > 0$, it holds that

$$\Pr[M = m \mid C = c] = \Pr[M = m].$$

- I.e., the distribution of M does not change conditioned on observing the ciphertext

Equivalent formulation

the **distribution of the ciphertext does not depend on the plaintext**, i.e., for any two messages $m, m^0 \in M$ the distribution of the ciphertext when m is encrypted should be identical to the distribution of the ciphertext when m^0 is encrypted. That is, for every $m, m^0 \in M$, and every $c \in C$, we have

$$\Pr[\text{Enc}_K(m) = c] = \Pr[\text{Enc}_K(m^0) = c]$$

it is **impossible to distinguish an encryption** of m from an encryption of m^0 , **since the distributions of the ciphertext are the same** in each case.

Perfect secrecy

- An encryption scheme $(\text{Gen}, \text{Enc}, \text{Dec})$ with message space M is perfectly secret if and only if Equation below holds for every $m, m^0 \in M$ and every $c \in C$.

$$\Pr[\text{Enc}_K(m) = c] = \Pr[\text{Enc}_K(m^0) = c]$$

- for any scheme, any distribution on M , any $m \in M$ for which $\Pr[M = m] > 0$, and any $c \in C$, we have

$$\begin{aligned}\Pr[C = c \mid M = m] &= \Pr[\text{Enc}_K(M) = c \mid M = m] \\ &= \Pr[\text{Enc}_K(m) = c \mid M = m] \\ &= \Pr[\text{Enc}_K(m) = c],\end{aligned}$$

$$\Pr[M = m \mid C = c] \cdot \Pr[C = c] = \Pr[C = c \mid M = m] \cdot \Pr[M = m]. \quad (2.3)$$

Perfect secrecy

$$\Pr[M = m \mid C = c] \cdot \Pr[C = c] = \Pr[C = c \mid M = m] \cdot \Pr[M = m].$$

$$\begin{aligned}\Pr[\text{Enc}_K(m) = c] &= \Pr[C = c \mid M = m] \\ &= \Pr[C = c] \\ &= \Pr[C = c \mid M = m'] = \Pr[\text{Enc}_K(m') = c]\end{aligned}$$

Perfect (adversarial) indistinguishability

- Adversary **passively observing a ciphertext** and then trying to **guess which of two possible messages** was encrypted.
- An adversary A first **specifies two arbitrary messages** $m_0, m_1 \in M$.
- A key k is generated using Gen.
- One of the two messages specified by A is chosen (each with probability $1/2$) and encrypted using k ; the resulting ciphertext is given to A.
- A outputs a “guess” as to which of the two messages was encrypted; A succeeds if it guesses correctly

Perfectly indistinguishable

The adversarial indistinguishability experiment $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}}$:

1. The adversary \mathcal{A} outputs a pair of messages $m_0, m_1 \in \mathcal{M}$.
2. A key k is generated using Gen , and a uniform bit $b \in \{0, 1\}$ is chosen. Ciphertext $c \leftarrow \text{Enc}_k(m_b)$ is computed and given to \mathcal{A} . We refer to c as the challenge ciphertext.
3. \mathcal{A} outputs a bit b' .
4. The output of the experiment is defined to be 1 if $b' = b$, and 0 otherwise. We write $\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}} = 1$ if the output of the experiment is 1 and in this case we say that \mathcal{A} succeeds.

it is trivial for \mathcal{A} to succeed with probability $1/2$ by outputting a random guess

Perfectly indistinguishable

- *Encryption scheme $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ with message space M is perfectly indistinguishable if for every A it holds that*

$$\Pr [\text{PrivK}_{\mathcal{A}, \Pi}^{\text{eav}} = 1] = \frac{1}{2} .$$

- *Encryption scheme Π is perfectly secret if and only if it is perfectly indistinguishable.*
- *Vigen`ere cipher is not perfectly indistinguishable*

One-time pad

- Patented in 1917 by Vernam
 - Recent historical research indicates it was invented (at least) 35 years earlier
- Proven perfectly secret by Shannon (1949)


One-time pad

- Let $\mathcal{M} = \{0,1\}^n$
- Gen: choose a uniform key $k \in \{0,1\}^n$
- $\text{Enc}_k(m) = k \oplus m$
- $\text{Dec}_k(c) = k \oplus c$

- Correctness:

$$\begin{aligned}\text{Dec}_k(\text{Enc}_k(m)) &= k \oplus (k \oplus m) \\ &= (k \oplus k) \oplus m = m\end{aligned}$$

One-time pad



Perfect secrecy of one-time pad

- *The one-time pad encryption scheme is perfectly secret.*

Fix an integer $\ell > 0$. The message space \mathcal{M} , key space \mathcal{K} , and ciphertext space \mathcal{C} are all equal to $\{0, 1\}^\ell$ (the set of all binary strings of length ℓ).

- Gen: the key-generation algorithm chooses a key from $\mathcal{K} = \{0, 1\}^\ell$ according to the uniform distribution (i.e., each of the 2^ℓ strings in the space is chosen as the key with probability exactly $2^{-\ell}$).
- Enc: given a key $k \in \{0, 1\}^\ell$ and a message $m \in \{0, 1\}^\ell$, the encryption algorithm outputs the ciphertext $c := k \oplus m$.
- Dec: given a key $k \in \{0, 1\}^\ell$ and a ciphertext $c \in \{0, 1\}^\ell$, the decryption algorithm outputs the message $m := k \oplus c$.

Secrecy of the one-time pad

PROOF We first compute $\Pr[C = c \mid M = m]$ for arbitrary $c \in \mathcal{C}$ and $m \in \mathcal{M}$ with $\Pr[M = m] > 0$. For the one-time pad, we have

$$\begin{aligned}\Pr[C = c \mid M = m] &= \Pr[K \oplus m = c \mid M = m] \\ &= \Pr[K = m \oplus c \mid M = m] \\ &= 2^{-\ell},\end{aligned}$$

Using the above result, we see that for any $c \in \mathcal{C}$ we have

$$\begin{aligned}\Pr[C = c] &= \sum_{m \in \mathcal{M}} \Pr[C = c \mid M = m] \cdot \Pr[M = m] \\ &= 2^{-\ell} \cdot \sum_{m \in \mathcal{M}} \Pr[M = m] \\ &= 2^{-\ell},\end{aligned}$$

Secrecy of the one-time pad

where the sum is over $m \in \mathcal{M}$ with $\Pr[M = m] \neq 0$. Bayes' Theorem gives:

$$\begin{aligned}\Pr[M = m \mid C = c] &= \frac{\Pr[C = c \mid M = m] \cdot \Pr[M = m]}{\Pr[C = c]} \\ &= \frac{2^{-\ell} \cdot \Pr[M = m]}{2^{-\ell}} \\ &= \Pr[M = m].\end{aligned}$$

We conclude that the one-time pad is perfectly secret.

Drawbacks of one time pad

- The key is as long as the message
- Is only secure if used once
- It is easy to see that encrypting more than one message with the same key leaks a lot of information.

$$c \oplus c^0 = (m \oplus k) \oplus (m^0 \oplus k) = m \oplus m^0$$

Limitations of Perfect Secrecy

- that any perfectly secret encryption scheme must have a key space that is at least as large as the message space

THEOREM 2.11 *If $(\text{Gen}, \text{Enc}, \text{Dec})$ is a perfectly secret encryption scheme with message space \mathcal{M} and key space \mathcal{K} , then $|\mathcal{K}| \geq |\mathcal{M}|$.*

Shannon's Theorem

THEOREM 2.12 (Shannon's theorem) *Let $(\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , for which $|\mathcal{M}| = |\mathcal{K}| = |\mathcal{C}|$. The scheme is perfectly secret if and only if:*

- 1. Every key $k \in \mathcal{K}$ is chosen with (equal) probability $1/|\mathcal{K}|$ by Gen .*
- 2. For every $m \in \mathcal{M}$ and every $c \in \mathcal{C}$, there is a unique key $k \in \mathcal{K}$ such that $\text{Enc}_k(m)$ outputs c .*

Summary

Discussed

- Probabilities
- Perfect Secrecy
- One-Time pad
- Shannon's Theorem on perfect secrecy